


Social Welfare Under Heterogeneous Time Preferences

Sarvin Bahmani
University of Liverpool

Soumyajit Paul, Sven Schewe,
Shadi Kalat, Ashutosh Trivedi.



Why Heterogeneous Time Preferences?

Key Question

What should an agent do when everyone agrees on the rewards, but disagrees about how much the future is worth?



Climate policy




Infrastructure



Investment

The Value of Delayed Rewards




A reward becomes **less valuable** when it is delayed.


Do People Value the Future Differently?




Yes! People value the future differently.



Low discounting
(impatient)



Moderate discounting



High discounting
(patient)

Classical Setting

Different principals may have different rewards but **same** discount factor.

Our Setting

Principals may also discount future rewards **differently**.

One Decision-Maker



Agent , Decision maker

Multiple Principals



Principals, stakeholders

Social Welfare Objective



Principal i 's payoff under policy π

$$V_i^\pi = \mathbb{E}_\pi \left[\sum_{t=0}^{\infty} \lambda_i^t R_i^t \right]$$


Social Welfare

$$SW(\pi) = \sum_i V_i^\pi$$

Optimisation Objective

$$\max_{\pi} SW(\pi)$$

Alice and Bob, Open a Restaurant!




Patient investor
Alice $\lambda : 2/3$



Impatient investor
Bob $\lambda : 1/3$

Social Welfare in different strategies



Policy 1: a^ω

individual payoff: $\frac{3}{1-\lambda}$

Policy 2: b^ω

individual payoff: $-1 + 6\frac{\lambda}{1-\lambda}$


Policy 3: $a^k b^\omega$

individual payoff: $3\frac{1-\lambda^k}{1-\lambda} - \lambda^k + 6\frac{\lambda^{k+1}}{1-\lambda}$

Social Welfare

$$SW(\pi) = \sum_i V_i^\pi$$

Social Welfare in different strategies



Policy 1: a^ω


individual payoff: $\frac{3}{1-\lambda}$

Policy 2: b^ω



individual payoff: $-1 + 6 \frac{\lambda}{1-\lambda}$

Policy 3: $a^k b^\omega$



individual payoff: $3 \frac{1-\lambda^k}{1-\lambda} - \lambda^k + 6 \frac{\lambda^{k+1}}{1-\lambda}$



Social Welfare in different strategies



Heterogeneous discounting turns time
itself into **Strategically relevant memory**.



How to get the best Social Welfare?

Negative Results

- Pure stationary strategies?
- Randomised strategies?

Simple stationary behavior is Hard

Positive Results



- ✓ Optimal strategies exist in a simple class.
- ✓ Pure finite-memory counting strategies

Structured finite memory is easy.

Long-run Behavior: Lexicographic Tail

Order principals by patience: $\lambda_0 > \lambda_1 > \lambda_2 > \lambda_3 > \dots > \lambda_n$


1. Optimize the MDP for the most patient principal.
 - a. Restrict to actions that are optimum.
2. Break ties using the next most patient principal.
3. Continue until last principals.



Deviation from Long Term Strategy

M_n
long-run tail

For




$$\sum_{s'} p(s' | s, \dots)$$

$$\lambda_i^t \Delta_0(s, a, i)$$

Social effect

$$\sum_i \Delta_t(s, a, i) = \sum_i \lambda_i^t \Delta_0(s, a, i)$$



After a finite horizon, deviations cannot help

M_n
long-run tail

for every deviation, there exists a finite horizon k such that, for all $t \geq k$:

$$\sum_i \Delta_t(s, a, i) \geq 0$$




① Why k is small under spacing?

A typical bound has the shape $k = O\left(\frac{\log(\text{value ratio})}{\log(\lambda_i/\lambda_{i+1})}\right)$


✓ Outcome

Only the prefix before K needs time-dependent optimization; after that, follow the stationary tail.

Scalability Trends



Synthesis Algorithm



Complexity Contrast

Insist on stationary simplicity	Allow structured finite memory
▪ Loses social welfare ▪ Randomization may help but does not solve the structural issue ▪ Stationary/positional threshold is NP-hard	✓ Attains optimal social welfare ✓ Pure finite counting strategies are sufficient ✓ Computable in polynomial time

- **Simple stationary behaviour is too restrictive**
 - ✓ **Structured finite memory is enough**

Takeaways

Thank you!

- 1 — Heterogeneous time preferences change the temporal structure of optimal control.
- 2 — Socially optimal behavior may need memory even when rewards are identical.
- 3 — Optimal strategies have a simple form $\pi^* = \underbrace{[\text{finite prefix}]}_{\text{counting memory}} \cdot \underbrace{[\pi_{\text{long}}^\omega]}_{\text{Stationary tail}}$
- 4 — Under reasonable spacing, synthesis is polynomial time.



Different notions of patience do not merely reweight rewards;
they change the structure of optimal behavior.


NP-Hardness

$$\phi = (x \vee \neg y \vee z) \wedge (\neg x \vee \neg y \vee z) \wedge (\neg x \vee y \vee \neg z)$$

$$\lambda_0 = 0.54$$

$$\lambda_1 = 0.40$$

$$\Sigma_i \left(0 + 0 \cdot \lambda_i + 0 \cdot \lambda_i^2 - \lambda_i^3 + \frac{\lambda_i^4}{1 - \lambda_i} \right)$$



$$\Sigma_i \left(0 + 0 \cdot \lambda_i + \lambda_i^2 - \frac{\lambda_i^3}{1 - \lambda_i} \right)$$